PRODUCTBEACON RESEARCH · CYBER RADAR

The Checkpoint Moved Inside the Model

Anthropic put data-loss prevention inside its own infrastructure. Four incumbent vendors plugged into it the same day.

THE LAUNCH

Anthropic built the inspection point, then published the protocol so anyone could plug in.

- On **August 5**, Anthropic launched **inference hooks** in beta for Claude Enterprise. Every prompt routes to a security server the customer runs, which returns an allow or deny verdict, and the model proceeds only once it has one.
- The same check runs on tool calls made through **Model Context Protocol** connectors, skills and plugins. The tool's response is inspected before it returns to the model.
- One organisation-level setting covers chat, Claude Code and Claude Cowork.
- The protocol is open and webhook-based with a published schema, so any security vendor can build to it.

WHO PLUGGED IN

Anthropic named four data-loss-prevention vendors. All four shipped that day. Others shipped uninvited.

- Anthropic's own post tells customers to point the hook at "**Netskope, Palo Alto Networks, Proofpoint, Zscaler** or an AI security server you built in-house."
- **Proofpoint** applied its existing policies. **Netskope** shipped One Data Loss Prevention On Demand, in private beta. **Zscaler** shipped Zero Trust enforcement. **Palo Alto Networks** connected Prisma AIRS to its Enterprise data-loss-prevention engine.
- Not named, shipped anyway: **Check Point**, **Cyera**, and newcomers **Reco** and **Akto**.
- Asked whether this cannibalises its own network products, Proofpoint said no, and that hooks add "coverage for usage that some network tools may not fully observe."

THE READ

For years the argument was where to stand. A model provider has now answered it for its own surface.

The security industry has spent a decade arguing whether data policy belongs on the network, on the endpoint or in the browser. Anthropic built the checkpoint into its own infrastructure and invited the incumbents to supply the verdict.

- Our read, and it is an interpretation: as far as we can tell, this is the strongest single week of evidence yet that a data-control layer is forming around the Model Context Protocol.
- It arrived in a shape we did not predict. The model provider owns the checkpoint. The security vendors supply the policy.
- One name is missing. **Cyberhaven** has promised a Claude Enterprise connector as “available in August 2026” since July 23, and it did not appear among the integrators.

THE MONEY

Then the deal machine stopped. One announcement, and no closes across eleven pending deals.

- The only new cyber-security acquisition was **AXA XL** agreeing on August 6 to buy the rest of London consultancy **S-RM**. Terms undisclosed, completion expected by end-September, subject to regulatory approval.
- No closes, clearances or terminations on any tracked pending deal. SecurityWeek’s deal feed shows nothing after July 30, and N2K CyberWire has published nothing since its August 5 issue.
- Disclosed funding fell to roughly **$205 million** across three rounds, down about 65% on the prior week.
- **Obsidian Security** raised $85 million at a $1.1 billion valuation. As far as we can tell, it is the first machine-identity vendor in this cycle to raise rather than sell.

THE CUT

Rapid7 cut about 12% of its workforce, and its recurring revenue is now going backwards.

- Revenue **$210.9 million, down 1.5%** year over year. Third-quarter recurring revenue guided down to about **$812 million, a 3% decline**.
- Approximately **12% of the workforce** was notified that positions were being eliminated, with $10 to $11 million in restructuring charges.
- Chief executive Wael Mohamed said customers “want us to go deeper in Detection and Response and Exposure Management, not wider.”
- Shares rose about 9%. Our read: the market priced the cost discipline and, we believe, looked past the shrinking revenue base, which is the more durable signal.

THE PRINTS

Two more prints landed the same day and went in opposite directions.

- **Cloudflare**: revenue **$696.1 million, up 36%**, full-year guidance raised, shares closed up 7.86%. Chief executive Matthew Prince called it “a fundamental rewrite of the Internet for machine-to-machine traffic.”
- **Akamai**: revenue up 5%, security up 10% and now its largest segment, cloud up 39%. But it reconfirmed revenue guidance while **cutting full-year profit guidance**, and GAAP operating margin fell 8 points.
- Same AI-infrastructure thesis, opposite near-term margin outcomes. As we read it, one is being paid for the build-out and the other is paying for it.

THE GATEKEEPER

And OpenAI decided who is allowed to sell frontier offensive capability.

- On **August 10** it launched **GPT-5.6-Cyber**, which it says is configured with a lower refusal rate for dual-use tasks, and split its Daybreak partner programme into two tiers: **Blue** for defensive frontier models, **Red** for the cyber-specific ones.
- Access runs only through approved partners, now including **Palo Alto Networks, CrowdStrike, Fortinet, Akamai, Cloudflare** and **Sophos**.
- Help Net Security's reading of the design: access to the models stays with the partner and is not transferred to the customer.
- Our read, and it is an interpretation: capability a vendor could once buy is becoming capability a vendor must be admitted to, and the admission list is itself becoming a competitive asset.

SOURCES

Every move here is public and dated.

The sources behind the moves in this carousel:

COMPANY AND PRIMARY ANNOUNCEMENTS

Anthropic · Proofpoint · Netskope · Zscaler · Palo Alto Networks · Check Point · Cyera · Reco · Akto · Cyberhaven · OpenAI · Rapid7 · Cloudflare · Akamai · Obsidian Security · AXA XL

BUSINESS AND TECHNOLOGY PRESS

SecurityWeek · Help Net Security · N2K CyberWire · The Next Web · Unite.AI · Security Boulevard · Tech Wire Asia · PR Newswire · Insurance Journal · StockTitan · Investing.com · Benzinga

Full citations, claim by claim and link checked, available on request.