
PRODUCTBEACON

ProductBeacon – State of Workforce Engagement Management 2026: The Agent- Assist & Real-Time Guidance Market

Yohay Etsion

Managing Director, ProductBeacon

2026-07-20

Disclosure: *This is ProductBeacon's independent, impartial market research, based solely on publicly verifiable signals, published on the open web. No vendor sponsors, no paywalled data, no analyst-firm reuse. It reflects ProductBeacon's independent view and is not investment advice. Full independence and sourcing methodology: [Methodology](#) →.*

By Yohay Etsion · Managing Director, ProductBeacon · former NICE product leader (Recording / Workforce Optimization), 2008–2021.

Former-employer independence: The author was employed by NICE (9 April 2008 to 20 May 2021), holding senior product-leadership responsibility for the Recording / Workforce-Optimization line of business – the very domain this report grades NICE on – and remains bound by a confidentiality and trade-secret undertaking from that employment; this report relies solely on published material and contains no information derived from that employment, and the author holds no NICE equity.

4.1 What This Market Is

What this market is. Agent-Assist & Real-Time Guidance is the software that changes the outcome of a conversation **while it is happening**. It covers real-time guidance and nudges, in-the-moment knowledge surfacing, next-best-action, live compliance prompts, and the after-call-work (auto-summary, disposition) that closes the interaction. It is the **live moment** of the WEM stack: where Quality & Interaction Analytics scores the interaction *after* and Performance & Engagement coaches the agent *later*, this market acts *during*. It also absorbs the buyer-visible, knowledge-in-the-flow half of **knowledge management (KM)**.

The buyer is the live moment. In practice the buyer is the supervisor or CX-ops leader who wants the agent to say the right thing, find the right answer, and finish the after-call work without a hold, on this call.

Where the boundaries fall – and this is the report's sharpest one.

- **vs Quality & Analytics:** Quality & Interaction Analytics is *retrospective* (score the completed interaction); this market is *real-time* (act during it). Same transcript, different time.
- **vs Performance & Engagement:** Performance & Engagement develops the agent *between* calls (coaching, gamification); this market guides them *on* the call.
- **vs Autonomous CX, the control-surface boundary:** This is the line the whole report turns on. **Assist returns a suggestion to a person who acts; autonomous is the resolution itself: the AI takes the action, and the human is out of the loop.** The discriminator is *who performs the resolving action*, not how capable the AI is, and **a human-approval gate is what keeps a system in Agent-Assist**. Vendors draw this line *themselves*, with two separately-named SKUs on one engine: NICE Copilot vs Autopilot, Cresta Agent Assist vs AI Agent, Genesys Copilot vs autonomous agents, and Five9 assist vs agents that “Reason, Decide, and Take Action”. That is the single-engine thesis made physical: **one engine, two SKUs, gated by whether a human is in the loop** (see the Winning & Losing section below and the Convergence chapter).

What it IS NOT. It is not post-interaction quality management (QM), which is Quality & Interaction Analytics. It is not between-call coaching or gamification, which is Performance & Engagement. It is not autonomous deflection or virtual agents that resolve without a human, which is Autonomous CX on the other side of the control-surface line. And it is not standalone enterprise knowledge management: only the in-the-flow, buyer-visible half lives here.

Three common buyer misconceptions. 1. *“Agent assist is live transcription.”* Transcription is table-stakes; the product is the *guidance* on top — what the agent should say or do next, grounded in something (top-performer behaviors, your own data). 2. *“A copilot and an autonomous agent are different products.”* Increasingly they are the **same engine** exposed as two SKUs separated only by a human-approval gate, which is why the assist vendor and the autonomous vendor are usually the same vendor. 3. *“More AI guidance is always better.”* The mature pitch is *calibrated* guidance and *trusted* guidance: Five9 explicitly suppresses prompts for seasoned agents “so as not to annoy them”, and explainability/evaluation is the 2026 enterprise gate.




Figure 4.1 – Where Agent-Assist sits: the live-call timeline on the X-axis (from during the call to closing the call), the control surface on the Y-axis, with the human-approval gate as the boundary to Autonomous CX. Quality and Agent-Assist are the

adjacent-time boundaries; only the in-the-flow half of knowledge management is in category.

4.2 Market Size & Trends

Category-scoping note. The market is sized as **real-time assist to a human agent**: live guidance, knowledge-in-the-flow, next-best-action, and auto-summary / after-call work (ACW). **In scope**: the copilot/guidance layer a human agent uses live. **Out of scope**: autonomous deflection / virtual agents that resolve without a human (Autonomous CX, the control-surface line), and post-interaction QM (Quality & Interaction Analytics). This is the **newest and least cleanly-sized** WEM-core market. Agent-assist is still being separated from the broader Contact Center as a Service (CCaaS) AI blend it was born inside.

Market size: blended and single-source-thin. Exactly one widely-circulated source, market.us, names “real-time agent assist” as its own category: ~\$4.4B in 2024 growing to \$124.6B by 2034 at ~40% CAGR.¹ **The market.us figure** is reported single-source with improbably-hot growth: carried here as a claim, not settled. The directional anchor is the better-corroborated “**call center AI**” **blend** from Grand View and Fortune, with labeled outliers.²³ Every figure in it is marked **blended**, since agent-assist is one line item inside it. There is no clean standalone dollar total addressable market (TAM) here.

Trends that move the buyer. (1) **Auto-summary / after-call-work automation is the measurable killer app.** The figures are vendor-claimed: Five9 ~40% of agent time⁴, Verint ~30 sec/call⁵. It is the wedge that proves ROI before autonomous resolution. (2) **Real-time RAG (retrieval-augmented generation) knowledge-surfacing**, the buyer-visible half of KM, becomes a default assist capability and thins standalone in-agent KM.⁶ (3) **Next-best-action whisper guidance** is grounded in top-performer behavior.⁷ (4) **The assist-to-autonomous gradient** is the structural story: Gartner’s “80% autonomous by 2029”⁸ pull runs against the counter-signal that a large share of agentic projects get canceled⁹. Both feed the Convergence falsification condition.

Compliance. Real-time processing puts customer data in the live loop, so the HIPAA/PCI/GDPR exposure from Quality & Interaction Analytics applies in real time. On top of that, the **EU AI Act’s prohibition on workplace emotion/sentiment recognition (Art. 5(1)(f)) bites real-time agent-sentiment scoring** specifically (the live version of the Agent-Assist emotion-AI wall — binding the EU deploying employer’s use, narrow medical/safety carve-outs aside), with full application Aug 2026.¹⁰

4.3 The Contenders

Eight vendors are graded on the 3-bucket rubric (cap 8). The three tiers group vendors by market weight: **Gravity** is the largest incumbents that pull the market, **Attention** is the credible mid-players worth watching, and **Wildcard** is the challenger or thin, early entrant. Each card carries a verbatim vendor pillar with its dated URL inline; latency/accuracy/outcome figures are quoted as *vendor claims*. AI-architecture verbs are publicly-evidenced positioning, not internals.




Figure 4.3 – The eight contenders on platform posture (single-engine suite ↔ standalone independent) against control-surface span (assist-dedicated ↔ owns both an assist and an autonomous SKU). Cresta and Balto are the independents that own an autonomous adjacency, so they sit HIGH despite being independent; Sprinklr is the public-company thin-surface hard-case. Positions are attributive reads of published material, not data-derived.

Gravity tier

1. **NICE: CXone Mpower Copilot for Agents (Enlighten Copilot)**. NICE's pillar is "AI copilot for customer service agents" / "Simplify agent work, amplify human impact," across its paired Copilot pages (nice.com).¹¹ The pages promise "Real-time AI assistance... proactively delivers knowledge and next best steps". They add "Knowledge when you need it... The right information appears exactly when

needed". Copilot is the live-moment surface of the same Enlighten engine that scores and coaches; its **Autopilot** is the autonomous SKU on the other side of the control-surface line. NICE is a Convergence P1 proxy.

2. Genesys: Cloud AI Copilot. Genesys pitches "**AI Copilot Agents for Smarter Service**" (genesys.com).¹² On the capability page, "**knowledge, insights, next-best action support, translation and summarization come standard**" in "**a common copilot agent interface**". The copilot is native to the platform that already routes the interaction; this market is an anchor card, since Genesys's primary home tilts assist/autonomous. Genesys is a P1 proxy.

3. Cresta: Agent Assist (the independent that earns a *second* card). Cresta positions gen-AI-native real-time guidance as a distinct, separately-positioned product, one of four named products: "**Unlock your team's true potential with real-time Agent Assist... real-time hints, reminders, and workflows tailored to each conversation.**" The pillar adds "**near-zero latency**" and source-backed knowledge (cresta.com).¹³ A named outlet puts the latency "under 200 ms". *AI positioning*: gen-AI-native. Assist "guides them... without interacting with the customer directly", the human-in-the-loop side of its own AI-Agent line. The scale signal, from vendor and named outlet, is "surpassed \$100M ARR in 2026" (claim, not fact). **Cresta is present in two markets, with a Quality card, this Agent-Assist card, and an Autonomous Play (AI Agent)**: the one vendor with two cards.

Attention tier

4. Balto: Real-Time Guidance Platform (the assist-primary independent). Balto's pillar is "**AI Powered Real-Time Guidance Platform**". The homepage promises "**Help every agent perform like your best with the most advanced agent assist,**" pulling answers from knowledge bases live, "no holds to look up information" (balto.ai).¹⁴ Balto is built assist-first, not bolted onto a QM / Contact Center as a Service core; its **Togo** voice-AI is the autonomous adjacency. *Funding (primary-source)*: **~\$52M total** per Balto's own Series B release (\$10M Series A from Sierra Ventures, \$37.5M Series B from Stripes in Aug 2021, plus seed).¹⁵ *A later ~\$5.7M round and a ~\$57.6M total appear aggregator-only and are unconfirmed at primary source.*

5. Uniphore: Real-Time Guidance Agent (Business AI Suite). Uniphore pitches "**AI agent assist for the customer journey**" (uniphore.com).¹⁶ The product page promises "**Real-time guidance, contextual answers, and intelligent automation**", plus after-call-work automation. It adds "**surface the right answers... no search required**". Uniphore is the heavily-funded independent at the Gravity boundary. Its verified-primary record is **Series E \$400M (Feb 2022, NEA), ~\$2.5B post-money at that round, >\$610M total**.¹⁷ **A Series F of \$260M announced 22 Oct 2025** followed, with investors reported to include NVIDIA, AMD, Snowflake, Databricks, and NEA. **The Series F figure** is reported at this refresh by a single named outlet and is unverified at primary.

Uniphore's last established valuation mark is the ~\$2.5B Series E (Feb 2022). Total raised comes to **>\$870M**: the primary-verified Series E plus the single-outlet-reported Series F. Broader outlet totals of ~\$886M–\$961M remain estimates. Uniphore is graded Attention on symmetric assist-surface-depth grading: assist is one module of a broad suite.

6. Level AI: AgentGPT. Level AI positions AgentGPT as a distinct gen-AI-native assist product, separate from its Quality Auto-QA card: **“a secure, omniscient generative AI system... trained on a client’s proprietary customer conversational data,”** giving **“instant access to documents, links, hints, and warnings while you’re in a conversation”** (thelevel.ai).¹⁸ Its wedge against generic copilots is “trained on your data”. *Total funding is reconciled on its Quality & Interaction Analytics card, not restated here.*

7. Five9: Genius AI Agent Assist (Contact Center as a Service suite assist anchor). It pitches **“real-time guidance cards... personalized to their competency level”**. The calibration is explicit: **“a seasoned agent might receive no real time guidance... whereas new hires will be pushed an exact talk track.”** It adds GPT-powered transcription and summarization that **“reduces... after-call work (ACW)”** (five9.com).¹⁹ Five9 is public (NASDAQ: FIVN); this is its anchor card, and it is a Play on Quality, Workforce Management, and Performance & Engagement. Five9 is a P1 proxy.

Wildcard tier

8. Sprinklr: Agent Copilot (the thin/secondary-surface hard-case). Sprinklr is a public customer-experience suite (NYSE: CXM) whose assist surface is real and dated but secondary to its core: **“GenAI-Powered Support for Agents”** with **“real-time nudges that improve metrics like first call resolution (FCR) and average handle time”** (Spring '26, BusinessWire 7 Apr 2026).²⁰ *Wildcard rationale:* no genuinely-thin seed/Series-A **human-agent-assist-primary** startup with a dated primary-source event cleared the viability bright line for a thin-surface card. The modern assist independents (Balto, Uniphore, Cresta, Level AI) are all well-funded, and the thin privates in the live-moment space are *autonomous-CX-primary*: Sierra, Decagon, and Wonderful sit in Autonomous CX, pre-Series-B, and are held to strategy opinion only there. Carrying a public company as the thin-surface hard-case avoids a viability stretch on an unsourceable private. **No viability claim is at stake (public company).**

Present but graded elsewhere: Observe.AI is Quality-primary with an Agent-Assist Play (its assist is secondary to QM, and the Quality-vs-Autonomous call resolved to Quality); Cogito is the assist capability of Thoma Bravo CX, which is not covered here; Cresta, resolved above, is the one two-card vendor.

4.4 What the Vendors Are Doing

Play 1: Assist becomes the live-moment module of one AI platform

NICE Copilot, Genesys’s native copilot, and Five9 Genius carry the play for the suites. The independents (Balto, Cresta, Uniphore, Level AI) defend on latency, behavioral-guidance depth, and trained-on-*your*-data. *This play feeds Convergence P1.*

Play 2: The live moment is rebuilt around next-best-action, not transcripts

Cresta pitches “behavioral guidance... grounded in the behaviors of top performers”¹³. NICE pitches “next best steps.” “Live transcription” commoditizes, and differentiation moves to behavioral grounding plus latency.

Play 3: Knowledge-in-the-flow folds into assist, collapsing standalone KM

Cresta ships Knowledge Agent²¹. NICE promises “knowledge when you need it”. Genesys makes “automatic knowledge” standard. Standalone in-agent KM thins, a Convergence input.

Play 4: Real-time assist becomes the on-ramp to autonomous CX

Cresta sells Agent Assist beside AI Agent; Balto sells assist beside Togo; Five9 runs assist inside the Genius agentic motion. The “guide the human” engine and the “be the agent” engine converge. That pulls Agent-Assist into Autonomous CX, where the seat question begins.

Play 5: “Trust the copilot” becomes the enterprise gate, via evaluation + explainability

Sprinklr pitches “Autonomous Evaluation... explainable logs”²⁰. Cresta pitches “Responsible AI.” The same evaluation lens that validates a copilot validates an autonomous agent.

4.5 Funding & Fault-Lines

Vendor	Round / status	Valuation	Investor / acquirer
NICE	Public (NASDAQ/TASE)	Public equity	—
Genesys	Private (Permira/Hellman & Friedman)	not isolated here	—
Cresta	\$125M Series D (Nov 2024); vendor "\$100M ARR 2026" (claim)	\$1.6B	WiL + QIA
Balto	~\$52M total (primary-source) — \$10M A / \$37.5M B (Stripes, Aug 2021) / seed ¹⁵ ; \$57.6M is aggregator-only, unconfirmed	Not disclosed	Sierra Ventures, Stripes
Uniphore	Series E \$400M (Feb 2022, NEA) ¹⁷ ; Series F \$260M (22 Oct 2025, single named-outlet, unverified at primary) ; >\$870M total (primary-verified E + single-outlet F)	~\$2.5B (Series-E mark, Feb 2022)	NEA (Series E); NVIDIA/AMD/Snowflake/Databricks reported (Series F); now an acquirer: ActionIQ/Infoworks ²²
Level AI	(on Quality card, not restated)	—	Adams Street
Five9	Public (NASDAQ: FIVN)	Public equity	—
Sprinklr	Public (NYSE: CXM)	Public equity	—



Figure 4.5 – The primary-source funding events in this market: Balto's Series B, Uniphore's Series E and Series F (the F single-outlet-reported at this refresh, unverified at primary, Oct 2025), Verint's acquisition of Cogito, and Cresta's Series D. Aggregator-only figures (Balto's "\$57.6M total", Cresta's "\$100M ARR") are excluded. Hedges travel verbatim onto the labels. No distress event was found for any contender, so none is plotted. Vertical is not to scale.

The funding story is the spine made visible. The well-funded AI-led assist money mostly sits with vendors that *card on Quality & Interaction Analytics* (Cresta \$1.6B, Observe.AI \$1.4B).

Assist is a line of a bigger AI platform, not a standalone raise. The one dedicated real-time-guidance pure-play (Balto) runs on pre-2022 capital. And the suites *absorbed* assist rather than out-raising it: Verint bought Cogito for \$38.2M.²³ Uniphore is the outlier, ~\$2.5B-valued at a mark set at its Feb 2022 Series E¹⁷, and now itself an acquirer. Uniphore's \$260M Series F was reported Oct 2025 by a single named outlet, unverified at primary.

Fault-lines, disclosed rather than guessed. Balto's "\$57.6M total" did not reconcile at primary source: Balto's own release says ~\$52M¹⁵; the extra round and total are aggregator-only (PitchBook and Tracxn diverge), so the report carries ~\$52M with the caveat. A **name-collision** exists: an unrelated "Balto" acquired by Fubo in 2020 is NOT the contact-center vendor. **Cogito's "\$45M" is an aggregator cumulative** (Series C \$37M in 2018 plus an extension), not a round.²⁴ No distress event was found for any contender, so the report makes no viability claim. The viability bright line applies: strategy opinion only, absent a primary-source event.

4.6 Winning & Losing

These are my read of where this market is heading: each is anchored to public evidence already cited above, framed as opinion, and stated as a falsifiable prediction with a dated checkpoint.

Pattern Claim 1 – The Control-Surface, Not the Engine

Observation. The same vendors sell real-time *assist* and *autonomous* CX as two separately-named SKUs on one platform, gated by whether a human acts. The pairs are explicit. NICE sells **Copilot** vs **Autopilot**¹¹. Cresta sells **Agent Assist** (“guides... without interacting with the customer directly”) vs **AI Agent**¹³. Genesys pairs **Copilot** with autonomous agents¹². Five9 pairs **Agent Assist** with agents that “Reason, Decide, and Take Action”¹⁹ (dated vendor pages, accessed 2026-06-16). **My read.** The Agent-Assist / Autonomous CX boundary is a **control-surface** boundary (who performs the resolving action), not an engine boundary – vendors run one engine and expose it as two products separated by a human-approval gate. This is the single-engine thesis made physical, and it is *also* the pricing hinge: the assist SKU keeps a human seat; the autonomous SKU removes it. **Conditional prediction.** If by **Q4 2026** at least **3** Gravity-tier vendors publicly position assist and autonomous as **one platform with a human-in-the-loop gate** (rather than two separate products), the control-surface thesis is confirmed. If instead they keep them architecturally separate, the single-engine collapse is weaker than the marketing implies. **Sources.** The dated vendor SKU pages quoted in the Contenders section (NICE Copilot/Autopilot, Cresta Agent Assist/AI Agent, Genesys, Five9; accessed 2026-06-16) and the control-surface boundary set out earlier in this chapter.




Figure 4.6a – One engine exposed as an assist SKU (Copilot) or an autonomous SKU (Autopilot), gated by whether a human is in the loop. The falsifiable test is boxed at the foot of the figure.

Pattern Claim 2 – Auto-Summary Is the Beachhead

Observation. After-call-work automation (auto-summary, disposition) is the market's measurable ROI wedge, and it ships as a *standard* assist capability across NICE, Genesys, Five9, and Uniphore. The wedge figures are vendor-claimed: Five9 ~40% of agent time⁴, Verint ~30 sec/call⁵. **My read.** After-call work (ACW) automation is the beachhead that proves assist ROI *before* the buyer trusts autonomous resolution — it is low-risk (post-call, human-reviewed), quantifiable, and universal. The vendor that wins ACW credibility earns the right to sell the rest of the copilot, then the autonomous SKU. **Conditional prediction.** If by **H1 2027** auto-summary/ACW automation is a standard, default-on capability across **all** Gravity assist products (not an add-on), it has commoditized as the table-stakes beachhead and differentiation has moved up to live guidance and autonomy. If it is still sold as a premium add-on, the ROI case is not yet settled. **Sources.** The Market Size & Trends section's after-call-work figures (vendor-claimed Five9 ~40% of agent time, Verint ~30 sec/call) and Plays 1 and 3 above.




Figure 4.6b – After-call-work automation is the low-risk, measurable ROI wedge the buyer trusts before live guidance and, last, autonomous resolution. The vendor-claimed proof points and the falsifiable test are carried in the figure.

Pattern Claim 3 – The Seat Question Starts at This Boundary

Observation. Assist keeps the human seat, because it guides a person; the *same vendors'* autonomous SKUs remove it (Cresta AI Agent, Balto Togo¹⁴, NICE Autopilot, Five9 agentic CX). Pricing is starting to follow: autonomous-CX vendors publicly list per-resolution/outcome pricing while assist stays per-seat. **My read.** The per-seat → per-outcome pricing migration the Convergence chapter tracks *physically begins at the Agent-Assist / Autonomous-CX control-surface line*: the moment the engine that guided the human starts replacing the human, the unit being priced stops being a seat. Agent-Assist is therefore the earliest observable site where the human seat count starts to fall. **Conditional prediction.** If by **mid-2027** at least **2** Gravity-tier vendors price their assist SKU per-seat but their autonomous SKU per-resolution/outcome on the same platform, the pricing hinge is confirmed at this boundary. If autonomous stays per-seat-priced, the seat-reset thesis is delayed. **Sources.** Play 4 above and the control-surface boundary set out earlier in this chapter; the per-resolution pricing evidence carried in the Autonomous CX chapter.




Figure 4.6c – At the control-surface line, assist is priced per seat while the same vendors' autonomous SKU is priced per resolution – the earliest observable site where the human seat count starts to fall. The falsifiable test is boxed at the foot of the figure.

Winners

These verdicts translate the tier placements and plays above; they are my opinion, not a scoreboard. **NICE, Genesys, Five9**. They own the live moment as one surface of a single-engine platform (their Gravity-tier cards above); assist is a module of the AI platform the buyer already runs, not a separate purchase. **Cresta**. It is the independent that earns a second card, on generative-AI depth and named-outlet “under 200 ms” latency; it is the clearest depth-and-latency answer to the suites. **Balto**. It is the assist-primary independent, built assist-first rather than added onto a QM / Contact Center as a Service core.

Watch

The control-surface line itself. The question is whether assist and autonomous stay one human-gated platform or fork into separate products (Pattern Claim 1's Q4 2026 checkpoint). This single question decides the report's single-engine thesis, and Agent-Assist is the earliest place it is observable. **Uniphore**. It is the heavily-funded scale outlier whose assist is one module of a broad suite: its ~\$2.5B mark was set at its Feb 2022 Series E, and Uniphore's \$260M Series F was reported Oct 2025 by a single named outlet, unverified at primary. The watch item is whether that capital converts into live-moment leadership.

Losers

No graded contender here carries a public distress or failure signal: the Funding & Fault-Lines section found no distress event for any of the eight, so I name none rather than manufacture one. The clearest *restructuring* evidence in this market is absorption, not failure. The standalone assist/emotion-AI capability folded into a suite: Cogito was absorbed into Verint, now Thoma Bravo CX, for \$38.2M.²³ And per three of the plays above, knowledge-in-the-flow is collapsing the standalone in-agent knowledge-management surface. I read that as consolidation of the category, not the death of a vendor.

4.7 The Road Ahead (H2 2026 watchlist)

1. **The control-surface line:** whether ≥ 3 Gravity vendors keep assist and autonomous on one human-gated platform (Pattern Claim 1's checkpoint); the single most important Convergence signal.
2. **Auto-summary/ACW** going default-on across all suites (Pattern Claim 2's checkpoint), the beachhead commoditizing.
3. **Per-seat (assist) vs per-resolution (autonomous) pricing** appearing on the same platform (Pattern Claim 3's checkpoint), the seat-reset's first observable site.
4. **Copilot evaluation/explainability** becoming the enterprise procurement gate (Sprinklr "Autonomous Evaluation"; the same lens validates autonomous agents).
5. **The assist-to-autonomous ramp:** whether the Gartner "80% autonomous by 2029" pull or the "agentic projects canceled" counter-signal proves out, the live falsification input to Convergence.

1. market.us, "Real-Time AI Agent Assist Market Size | CAGR of 39%" — reports USD 4.4B (2024) to USD 124.6B by 2034 at 39.7% CAGR; single second-tier syndicated source (flagged). Accessed 2026-06-16. <https://market.us/report/real-time-ai-agent-assist-market/> ↩
2. Grand View Research, "Call Center AI Market Size & Share | Industry Report, 2030" — USD 1.99B (2024) to USD 7.08B by 2030 at 23.8% CAGR; blended call-center-AI envelope. Accessed 2026-06-16. <https://www.grandviewresearch.com/industry-analysis/call-center-artificial-intelligence-market-report> ↩
3. Fortune Business Insights, "Call Center AI Market Size, Share, Growth | Global Report [2034]" — USD 2.98B (2026) to USD 13.52B by 2034 at 20.8% CAGR; blended call-center-AI envelope. Accessed 2026-06-16. <https://www.fortunebusinessinsights.com/call-center-ai-market-109249> ↩
4. Five9, "Want to Increase Agent Productivity? Start with Automated Call Summarization" — vendor states AI Summaries saves up to 40% of an agent's time (named customer TruConnect cut ACW 40%). Accessed 2026-06-16. <https://www.five9.com/blog/want-increase-agent-productivity-start-automated-call-summarization> ↩↩
5. Capacity, "What Is Real-Time Agent Assist for Contact Centers? (2026 Guide)" — reports Verint ~30 sec/call ACW reduction across 30,000 agents; 20–40% handle-time reduction. Accessed 2026-06-16. <https://capacity.com/blog/real-time-agent-assist-for-contact-centers/> ↩↩
6. Capacity, "AI Agent Assist Software: Use Cases & Top 7 Tools (2026)" — real-time knowledge base delivers the relevant answer instantly on detected customer intent. Accessed 2026-06-16. <https://capacity.com/blog/ai-agent-assist-software/> ↩

7. CX Foundation, "Agent Assist in 2026: A Definition, Use Cases, & Best Practices" — real-time knowledge, next-best-action, and wrap-up as the named assist use cases. Accessed 2026-06-16. <https://cxfoundation.com/blog/agent-assist> ↩
8. Gartner, "Gartner Predicts Agentic AI Will Autonomously Resolve 80% of Common Customer Service Issues Without Human Intervention by 2029," 2025-03-05. Accessed 2026-06-16. <https://www.gartner.com/en/newsroom/press-releases/2025-03-05-gartner-predicts-agentic-ai-will-autonomously-resolve-80-percent-of-common-customer-service-issues-without-human-intervention-by-2029> ↩
9. Gartner, "2026 Hype Cycle for Agentic AI" — Gartner predicts 40%+ of agentic-AI projects canceled by end-2027 on cost, unclear value, and risk-control grounds. Accessed 2026-06-16. <https://www.gartner.com/en/articles/hype-cycle-for-agentic-ai> ↩
10. State of Surveillance, "The EU AI Act Takes Full Effect in August. Here's What It Actually Bans" — Art. 5(1)(f) workplace emotion-recognition prohibition (in force 2 Feb 2025); full application 2 Aug 2026. Accessed 2026-06-16. <https://stateofsurveillance.org/news/eu-ai-act-august-2026-biometric-surveillance-explainer/> ↩
11. NICE, "Copilot for Agents" and "Enlighten Copilot" product pages — vendor states "AI copilot for customer service agents," "proactively delivers knowledge and next best steps," "The right information appears exactly when needed." Accessed 2026-06-16. <https://www.nice.com/products/copilot-for-agents> and <https://www.nice.com/products/enlighten-copilot> ↩↩
12. Genesys, "Agent Assist / AI Copilot Agents" capability page — vendor states "knowledge, insights, next-best action support, translation and summarization come standard" in "a common copilot agent interface." Accessed 2026-06-16. <https://www.genesys.com/capabilities/agent-assist> ↩↩
13. Cresta, "Agent Assist" product page — vendor states "Unlock your team's true potential with real-time Agent Assist... real-time hints, reminders, and workflows tailored to each conversation" and "near-zero latency." Accessed 2026-06-16. <https://cresta.com/agent-assist> ↩↩↩
14. Balto, homepage — vendor states "AI Powered Real-Time Guidance Platform," "Help every agent perform like your best with the most advanced agent assist," "no holds to look up information"; Togo voice-AI positioned as its autonomous adjacency. Accessed 2026-06-16. <https://www.balto.ai/> ↩↩
15. Balto, "Balto Raises \$37.5M, Pushes to Close Execution Gap for Contact Centers" (Series B led by Stripes; stated total \$52M), 19 Aug 2021; and PR Newswire, "Conversation Intelligence Leader Balto Secures \$10 Million Series A Funding" (Sierra Ventures lead), 15 Oct 2020. Accessed 2026-06-16. <https://www.balto.ai/blog/balto-raises-37-5m-pushes-to-close-execution-gap-for-contact-centers/> and <https://www.prnewswire.com/news-releases/conversation-intelligence-leader-balto-secures-10-million-series-a-funding-301152942.html> ↩↩↩
16. Uniphore, "Real-Time Guidance Agent" product page — vendor states "AI agent assist for the customer journey," "Real-time guidance, contextual answers, and intelligent automation," "surface the right answers... no search required." Accessed 2026-06-16. <https://www.uniphore.com/products/customer-service-ai/real-time-guidance-agent/> ↩
17. Uniphore, "Uniphore Raises \$400 Million to Support Explosive Growth and Global Demand for Automating Conversations Across the Enterprise" (Series E, NEA lead, \$2.5B post-money, >\$610M total raised), 16 Feb 2022. Accessed 2026-06-16. <https://www.uniphore.com/press-releases/uniphore-raises-400-million-to-support-explosive-growth-and-global-demand-for-automating-conversations-across-the-enterprise/> ↩↩↩
18. Level AI, "AgentGPT" (thelevel.ai/agents) and BusinessWire, "Level AI Introduces the Future of Customer Service with First-of-its-Kind Generative AI Solution AgentGPT," 22 Feb 2023 — vendor states "a secure, omniscient generative AI system... trained on a client's proprietary customer conversational data." Accessed 2026-06-16. <https://thelevel.ai/agents/> and <https://www.businesswire.com/news/home/20230222006068/en/Level-AI-Introduces-the-Future-of-Customer-Service-with-First-of-its-Kind-Generative-AI-Solution-AgentGPT> ↩

19. Five9, "Agent Assist" capability page — vendor states "real-time guidance cards... personalized to their competency level," "a seasoned agent might receive no real time guidance... whereas new hires will be pushed an exact talk track," and transcription/summarization that "reduces... after-call work (ACW)." Accessed 2026-06-16. <https://www.five9.com/products/capabilities/agent-assist> ↩↩
20. Sprinklr, "Agent Copilot" product page and BusinessWire, "Sprinklr Unveils Next Wave of AI-Native Customer Experience Innovation with Spring '26 (26.4) Release," 7 Apr 2026 — vendor states "GenAI-Powered Support for Agents" with "real-time nudges that improve metrics like first call resolution (FCR) and average handle time." Accessed 2026-06-16. <https://www.sprinklr.com/products/customer-service/agent-copilot/> and <https://www.businesswire.com/news/home/20260407515870/en/Sprinklr-Unveils-Next-Wave-of-AINative-Customer-Experience-Innovation-with-Spring-26-26.4-Release> ↩↩
21. Cresta, "Cresta Launches Knowledge Agent for Contact Centers" (CMSWire) — Knowledge Agent surfaces real-time answers where agents work, atop Agent Assist. Accessed 2026-06-16. <https://www.cmswire.com/contact-center/cresta-launches-knowledge-agent-for-contact-centers/> ↩
22. Business Wire, "Uniphore Acquiring ActionIQ and Infoworks to Deliver the Industry's First Zero Data AI Cloud," 5 Dec 2024 (terms undisclosed). Accessed 2026-06-16. <https://www.businesswire.com/news/home/20241205133623/en/Uniphore-Acquiring-ActionIQ-and-Infoworks-to-Deliver-the-Industrys-First-Zero-Data-AI-Cloud> ↩
23. Verint, SEC Form 10-Q (FY2024) — Cogito acquisition \$38.19M cash at close (plus up to ~\$23.2M contingent), closed 11 Oct 2024; Verint → Thoma Bravo CX on the 26 Nov 2025 take-private close. Accessed 2026-06-16. <https://www.sec.gov/Archives/edgar/data/0001166388/000116638824000176/vrnt-20241031.htm> ↩↩
24. Business Wire, "Cogito Raises \$37M in Growth Equity Funding" (Series C, Goldman Sachs Growth Equity lead), 24 Jul 2018. Accessed 2026-06-16. <https://www.businesswire.com/news/home/20180724005244/en/Cogito-Raises-37M-in-Growth-Equity-Funding> ↩

Disclosures

DISCLOSURE.

This is ProductBeacon's independent, impartial market research, based solely on publicly verifiable signals, published on the open web. No vendor sponsors, no paywalled data, no analyst-firm reuse. It reflects ProductBeacon's independent view and is not investment advice. Full independence and sourcing methodology: [Methodology](#) →.

NOT ADVICE.

This report does not constitute investment, legal, tax, or accounting advice. No claim should be relied upon as the basis for any investment decision. The author holds no trading position in any named public security and is not compensated by any named vendor. Readers who use this report in investment contexts bear sole responsibility for their decisions.