PRODUCTBEACON

ProductBeacon — State of Workforce Engagement Management 2026 — The Convergence Synthesis — One Engine, One Seat-Count Reset

Yohay Etsion

Managing Director, ProductBeacon

2026-07-20

Disclosure: *This is ProductBeacon's independent, impartial market research, based solely on publicly verifiable signals, published on the open web. No vendor sponsors, no paywalled data, no analyst-firm reuse. It reflects ProductBeacon's independent view and is not investment advice. Full independence and sourcing methodology: Methodology →.*

By Yohay Etsion · Managing Director, ProductBeacon · former NICE product leader (Recording / Workforce Optimization), 2008–2021.

Former-employer independence: The author was employed by NICE (9 April 2008 to 20 May 2021), holding senior product-leadership responsibility for the Recording / Workforce-Optimization line of business — the very domain this report grades NICE on — and remains bound by a confidentiality and trade-secret undertaking from that employment; this report relies solely on published material and contains no information derived from that employment, and the author holds no NICE equity.

The Bet

The thesis the whole report is built to test, in one sentence:

> **As WEM buyers renew through 2026–2027, the classic WFO stack — recording → quality → workforce management → coaching — is consolidating onto a single AI brain, and the value WFO captures is migrating from per-seat licensing toward per-outcome / per-resolution pricing as autonomous agents absorb the interactions a human seat used to handle.**

It is a *bet*, not a description: falsifiable, with a dated falsification condition (the unified verdict below). It rests on three proxies. The five preceding markets were each structured to test one of them on **public, reader-checkable evidence**, never on a forecast:

- **P1: one engine** (the stack collapses onto a single AI brain): tested at the *product surface* (Proxy P1 below).
- **P2: seat → outcome** (the pricing unit changes): tested on *published price lists today* (Proxy P2 below).
- **P3: the seat absorbs** (autonomous agents take the volume a seat used to handle): tested on *deflection/fulfillment signals* — and this is the contested one (Proxy P3 below).

The honest headline, delivered before the evidence: **P1 is strongly evidenced and P2 is observable fact today; P3, the actual seat elimination, is a contested forecast, and it is the report's live falsification risk.** The repricing is real; the displacement is a bet.

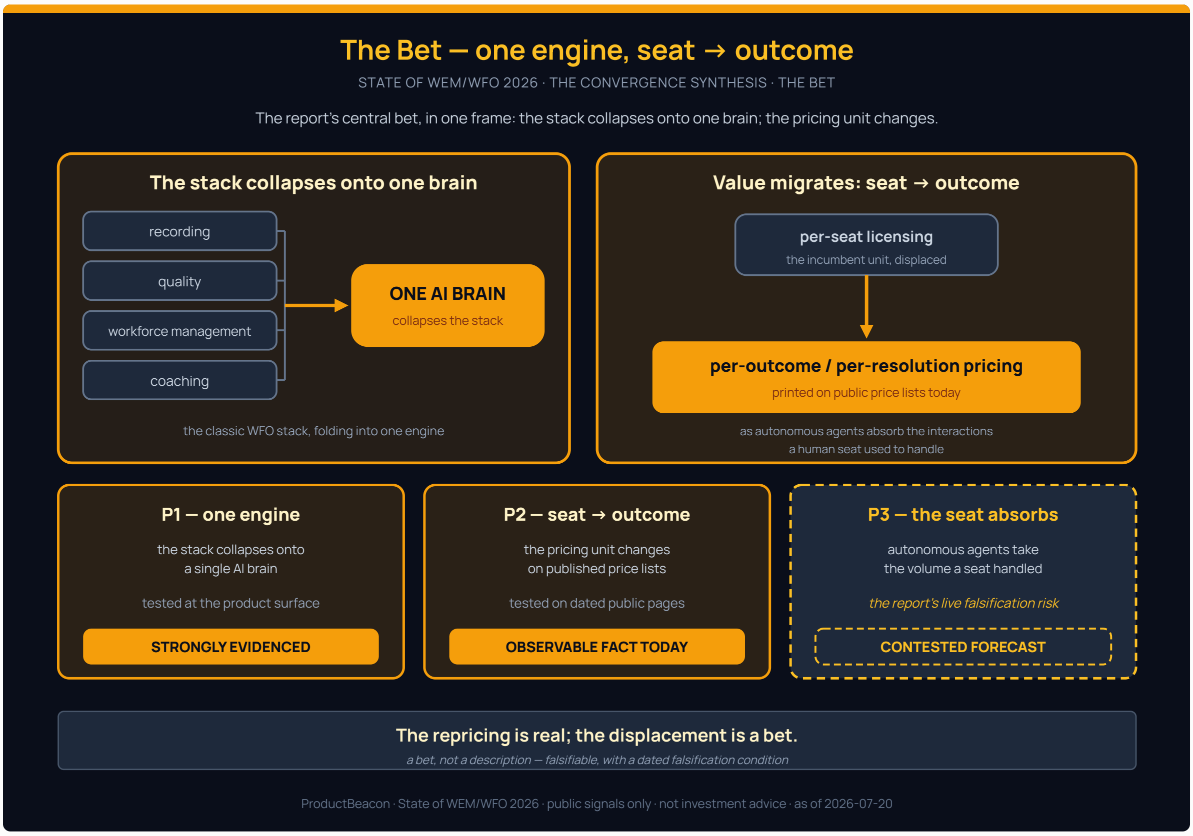

Figure C.1 — The bet in one frame: the classic stack collapses onto one AI brain, and value migrates from per-seat licensing toward per-outcome pricing. P1 and P2 are evidenced; P3 is the contested forecast.

Proxy P1: One engine, made physical

The single-engine claim is not an architecture diagram we drew; it is a **SKU boundary the vendors drew themselves**, and it is visible three ways across the markets:

1. **The multi-market card pattern.** The same suites card as *distinct, separately-positioned products* on market after market. **NICE carries a card on all five WEM-module markets** (Quality, WFM, Performance, Assist, Autonomous); Genesys carries one on five; Thoma Bravo CX (Verint+Calabrio) on three. Thoma Bravo CX is a *Play*, not a card, on Assist and Autonomous, where its surface is the absorbed Cogito. That breadth is not double-counting, since each card is a genuinely distinct product per the cross-market ledger. It is the single vendor positioning one platform across the whole stack. On 22 June 2026, after the research cut, Salesforce extended the same pattern beyond its Autonomous CX card with **Agentforce Contact Center WEM**, “the only WEM solution with AI, CRM, and every channel built in,” per the vendor page (salesforce.com). It adds workforce management and quality management, with visibility of both human and AI activity.

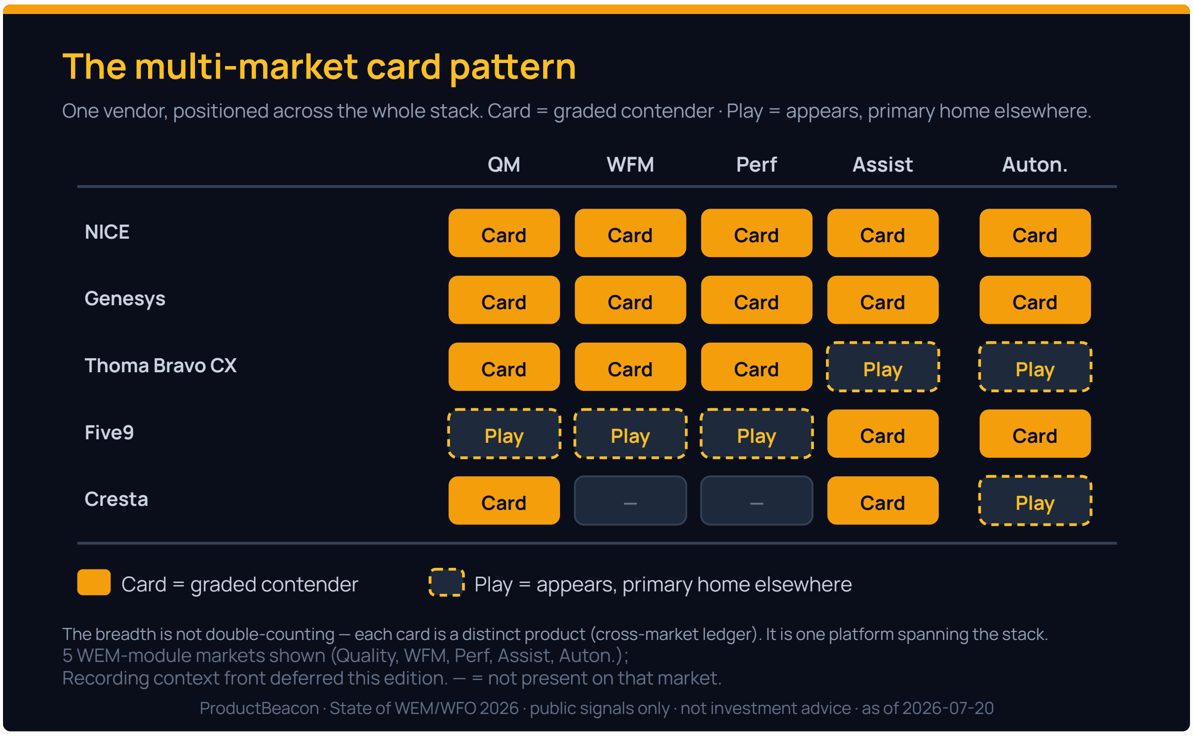

Figure C.2a — The multi-market card pattern: one vendor positioned across the whole WEM stack. Card = graded contender; Play = appears, primary home elsewhere.

2. **The named engine.** The suites increasingly give the engine *one name* spanning the modules. NICE's **Enlighten** scores in Quality, forecasts in WFM, coaches in Performance, guides in Assist, and resolves in Autonomous. Five9's **"Agentic Experience Engine"** explicitly powers both Agent Assist and the autonomous AI Agents.[1] One engine, many surfaces.

Freshness (to 20 Jul 2026): the naming surface moved on three vendors. **Verint**, merged with Calabrio and now presenting unified under the Verint brand, launched **Agent Factory**, a named orchestration environment for human + AI agent workforces beside its Workforce / Desktop / Quality Intelligence agentic capabilities (verint.com / BusinessWire, 23 Jun 2026). The counterweight is the CEO's own framing: "in a lot of ways… more repositioning of technology that we've had in place" (Dave Rhodes, attributed, CMSWire). **NICE**'s June 2026 re-architecture names NiCE Cognigy's agentic AI as "the brain of the platform," with three new layers: NiCE AI Agents, Agentic Engagement Plane, Agentic Analytics (nice.com, 9 Jun 2026). The spanning-engine NAME is migrating from Enlighten-only toward Cognigy-plus-Enlighten even as the one-platform story deepens. **Five9**'s Agentic Experience Engine name remains present across its public surfaces as observed this cycle (the launch release, current blog and Genius AI resource pages, and localized product pages), so Five9 counts in the Q4-2026 tally as-is. The observed drift is de-emphasis: the mid-2026 EN product page and the June 2026 Voice AI Agents release keep the name out of primary copy. That is a watch item, not fragmentation.

3. **The control-surface SKU gate (the physical proof).** The cleanest evidence is the Agent-Assist / Autonomous-CX boundary established in Agent-Assist and completed in Autonomous CX: the same engine is exposed as *assist* or *autonomous* depending only on **whether a human is in the loop**. NICE's own copy is the report's single clearest instance: its Mpower agents "support employees via

> **Mpower Copilot** or serve customers directly via **Mpower Autopilot**, tapping into the full CXone Mpower ecosystem."[2] The result is one platform, two SKUs, gated by who performs the action. Cresta (Agent Assist ↔ AI Agent) and Five9 (Agent Assist ↔ AI Agents on one engine) draw the same gate.

The per-market Winning/Losing sections kept returning to the same strategic consequence: the suite's autonomous (or assist, or coaching) SKU is a *configuration of the engine the buyer already runs*. That is the "one AI brain collapses the stack" claim rendered as a procurement reality rather than a slide.

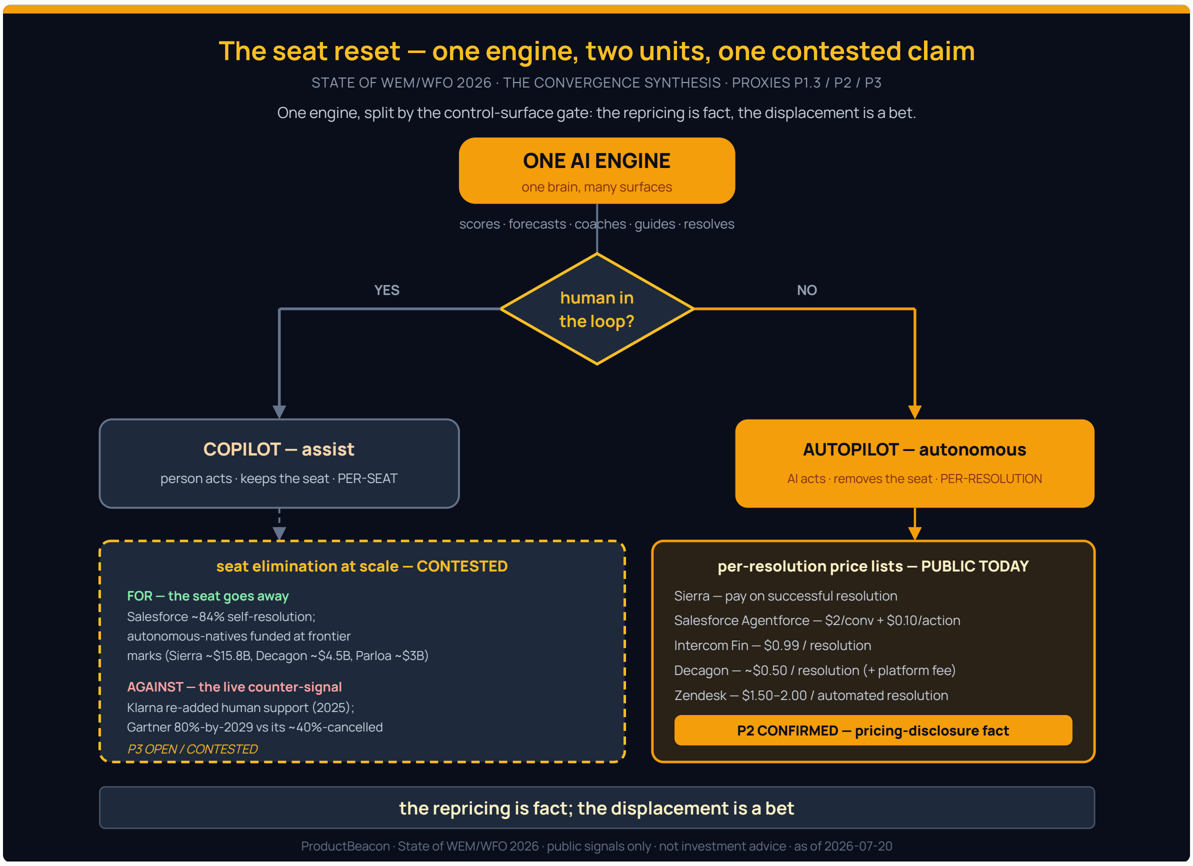

Figure C.2b — The seat reset: one engine, split by the control-surface gate into per-seat assist and per-resolution autonomous. The repricing (P2) is an observable fact; the displacement (P3) is contested.

Proxy P2: Seat → outcome, observable in price lists today

P2 is the proxy that most cleanly survives public testing, because, unlike the displacement forecast, the **pricing migration is already printed on dated public pages.** Autonomous CX is where per-seat licensing visibly breaks:

Vendor	Published unit	(market)
Sierra	outcome-based — “pay on successful resolution,” no charge on escalation[3]; extended to Horizon agents (“you pay for results, not tokens,” 16 Jul 2026)	Autonomous
Salesforce Agentforce	$2 / conversation + $0.10 / action (Flex Credits)[4]; pay-per-resolution announced for Help Agent (25 Jun 2026, GA July 2026 — $2 reported, not yet printed on the pricing page)	Autonomous
Fin (formerly Intercom)	$0.99 / outcome[5] — Salesforce acquisition announced 15 Jun 2026 (~$3.6B; close expected Salesforce fiscal Q4 2027)	Autonomous (adjacent)
Decagon	~$0.50 / resolution (+ platform fee) — reported; Decagon publishes no pricing page[6]	Autonomous
Zendesk	per automated resolution — construct printed on its pricing page; the $1.50–2.00 range as previously reported[7]	Autonomous (adjacent)

The unit of value is moving from the **seat** to the **resolved interaction**, and it is moving in public. The chain that makes this a WEM-wide finding, not just an Autonomous CX one, runs backward through the stack. WFM already re-described its planning unit as “humans + BPO + **AI agents**”: you cannot price a bot per seat. Agent-Assist located the pricing hinge at the control-surface line: assist keeps the seat, autonomous removes it. Autonomous CX is where the per-resolution price list actually appears. **P2 is confirmed as a pricing-disclosure fact:** reported as exactly that, never as a seat-count volume forecast. Freshness (2026-07-20): every anchor above was re-verified live or at Tier-1 corroboration; **none withdrawn, none reverted**. Salesforce announced a fourth pricing construct (pay-per-resolution for Help Agent, 25 Jun 2026). It also completed its acquisition of **m3ter**, a consumption-metering platform (announced 8 Jun 2026; completed 1 Jul 2026). m3ter is the billing plumbing for outcome pricing at platform scale.

Proxy P3: The contested displacement (where the bet is genuinely at risk)

P3 is the proxy the report refuses to overclaim. The *capability* direction is clear and advancing. Vendors have redefined the metric from **deflection** (“keep it from a human”) to **fulfillment** (“the agent executes the resolution end-to-end”): NICE with “AI that gets things done,”[8] Genesys with “execution,”[9] and Lorikeet with its “40-step refund flow.”[10] But the *displacement* claim, that autonomous agents remove human seats at scale, is genuinely two-sided on public evidence:

- **For:** Salesforce reports ~84% self-resolution on its own support[11]; the autonomous-native independents are funded at frontier marks (Sierra ~$15.8B[12], Decagon ~$4.5B[13], Parloa ~$3B[14]) on the premise that the seat goes away.

- **Against (the live counter-signal):** Klarna publicly said it had leaned too far into AI automation and moved to re-add human support (CEO Sebastian Siemiatkowski, public statements, 2025).[15] Gartner's "80% autonomous by 2029"[16] pull sits against its own projection that a large share of agentic projects are cancelled by 2027[17]. In-window (July 2026) reporting enriches both legs of the Klarna datapoint: the assistant's workload is reported grown to ~850 agent-equivalents (Forbes, 16 Jul 2026), while the re-added human support is reported as gig-structured — "an Uber type of set-up," per Siemiatkowski (The Guardian, 18 Jun 2026). The counter-signal stands: more textured, neither retracted nor repeated elsewhere.

The repricing is happening (P2, fact), fulfillment is rising (P3 capability, evidenced), but **seat elimination at scale is a forecast with a live, public counter-example.** The report stakes its credibility on keeping that distinction visible at every turn: pricing is fact, displacement is a bet.

The consolidation engines (the M&A spine across all markets)

The cross-market M&A pattern resolves into **three** structurally distinct consolidation engines, all pointing the same direction (fewer independent WEM vendors, more capability folded into multi-market platforms), by opposite mechanisms:

1. **NICE: organic / acqui-suite.** It is a public company building the single-engine on its own balance sheet and *buying capability into the owned stack*: Cognigy (~$955M, agentic → Autonomous)[18], LiveVox (~$350M, outbound)[21], Playvox (WEM/engagement)[19]. NICE integrates to own the product direction. The engine investment's reported line: **Q1 2026 AI ARR up 66% year over year** (NiCE Q1 2026 earnings release, May 2026). The 66% figure is the AI ARR line specifically, not total revenue: total revenue grew ~10%.

2. **Thoma Bravo CX: PE roll-up.** It is private capital *aggregating* best-of-breed into one platform: Verint ($2.0B take-private, closed 26 Nov 2025)[20] merged with Calabrio (which had itself absorbed Echo AI), plus Verint's earlier acquisition of Cogito (emotion-AI). Thoma Bravo CX aggregates to own the installed base. On WFM, this engine literally de-dups two top-tier WFM stacks into one.

3. **The autonomous-native venture wave.** This is a *third* engine the prior State-of-X reports did not have: frontier-funded independents (Sierra, Decagon, Parloa) raising at marks that let them compete with the suites for the *category*, not just a feature. This is the one market (Autonomous CX) where independent capital is genuinely contesting the platform.

The freshness window (to 20 Jul 2026) added buy-side datapoints in the same direction. **Salesforce** signed a definitive agreement to acquire **Fin (formerly Intercom)**, the per-resolution-priced independent anchor, for approximately **$3.6B** (announced 15 Jun 2026). The deal is announced, not closed: the close is expected in Salesforce's fiscal Q4 2027, subject to regulatory clearances. Salesforce also **completed** its acquisition of **m3ter** (consumption metering; 1 Jul 2026). **Genesys** acquired **Pinkfish** (agentic orchestration; 30 Jun 2026, terms undisclosed), and **Sprinklr** acquired the assets of **ViralMoment** (social-video intelligence; 28 May 2026, terms undisclosed). The direction is unchanged: fewer independents, more capability folded into multi-market platforms.

The first two integrate/aggregate the installed base; the third tries to out-run it. The falsification condition (the Winning & Losing section below) is, in effect, a bet on which engine wins net-new.

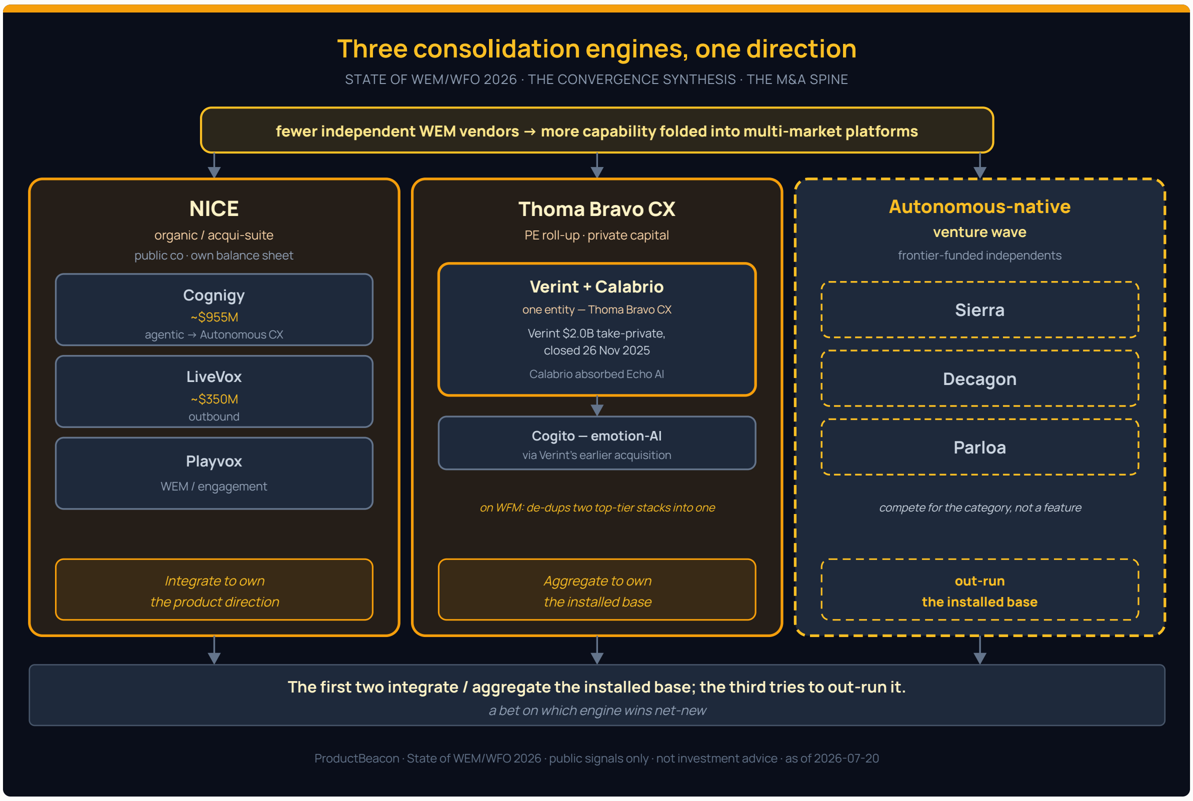

Figure C.5 — Three consolidation engines, one direction: NICE integrates, Thoma Bravo CX (Verint + Calabrio) aggregates, and the autonomous-native venture wave tries to out-run the installed base.

Winning & Losing: the unified verdict

This is my read across the markets: each claim is anchored to public evidence already cited in the chapters above, framed as opinion, and stated as a falsifiable prediction with a dated checkpoint.

Cross-Market Pattern Claims

Pattern Claim 1 — The One Engine Is Real at the Product Surface (P1 verdict)

Observation. The same suites card as distinct products across the full WEM stack: NICE on all five WEM-module markets, Genesys on five, Thoma Bravo CX on three. They do it increasingly under one named engine (NICE Enlighten; Five9 "Agentic Experience Engine"). And they expose assist vs autonomous as one platform gated by a human-in-the-loop SKU (NICE Mpower Copilot/Autopilot). All of it is on dated public surfaces cited in the five graded-market chapters. **My read.** The "one AI brain collapses the stack" claim is confirmed *at the product/SKU surface*. It is not proof of one internal codebase, which no public source can establish, and which the report never asserts. It is the way the category leaders now position and price. The collapse is a commercial reality before it is necessarily an architectural one. **Conditional prediction.** If by **Q4 2026** at least **3** Gravity-tier vendors publicly name one engine/platform spanning **≥3** WEM modules *and* a Copilot/Autopilot-style human-gated SKU pair, P1 is confirmed at the surface. If the modules stay separately-branded, separately-priced products, the single-engine story is weaker than the marketing. **Sources.** The five graded-market contender cards; the cross-market vendor-card pattern set out in this section; the single-engine-gated-SKU evidence in the Autonomous CX chapter.

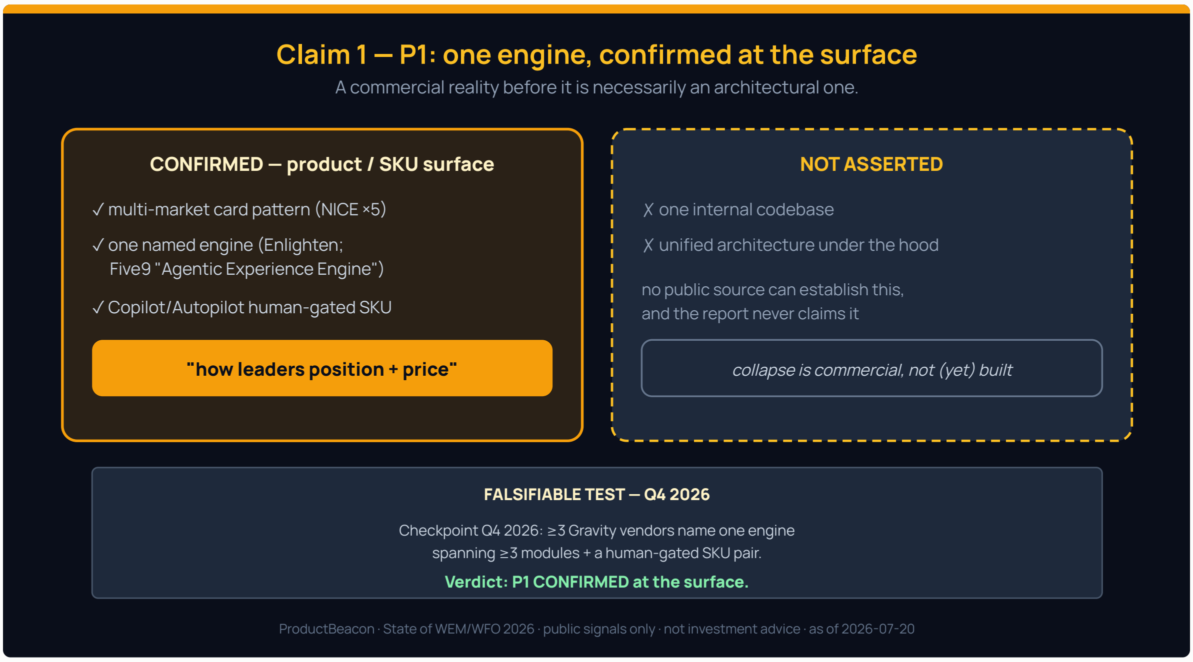

Figure C.6a — Pattern Claim 1 (P1): the one engine is confirmed at the product and SKU surface, not asserted as one internal codebase.

Pattern Claim 2 — The Repricing Is Real; the Displacement Is Not Yet (P2 confirmed / P3 contested)

Observation. Per-resolution / per-outcome pricing is published on dated public pages today. Sierra is outcome-based; Agentforce is $2/conversation + $0.10/action. Fin (formerly Intercom) is $0.99/outcome; Decagon is ~$0.50 (reported). Zendesk prices per automated resolution. The anchor pages were re-verified 2026-07-20; none withdrawn or reverted. The seat-displacement evidence, by contrast, is two-sided: Salesforce ~84% self-resolution vs Klarna's public reversal; Gartner 80%-by-2029 vs its ~40%-cancelled. **My read.** The seat → outcome **repricing** is an accomplished, observable fact; the seat **elimination** is a contested forecast. The report's central, defensible finding is precisely this asymmetry. The market has already changed how it *prices* WFO faster than it has changed how many humans it *staffs*. The pricing leads the displacement, and may lead it for years. **Conditional prediction.** If by **H1 2027** at least **2** Gravity-tier vendors price a flagship customer-facing CX-agent SKU primarily off per-seat (to per-resolution/conversation/action) on a dated public page, P2 is confirmed. The displacement (P3) remains *open* until sustained (non-pilot) autonomous fulfillment above **50%** on a primary channel is publicly reported without a Klarna-style reversal. **Sources.** The Autonomous CX pricing evidence; the seat-question chain traced through the WFM and Agent-Assist chapters; the Autonomous CX Terrain section.

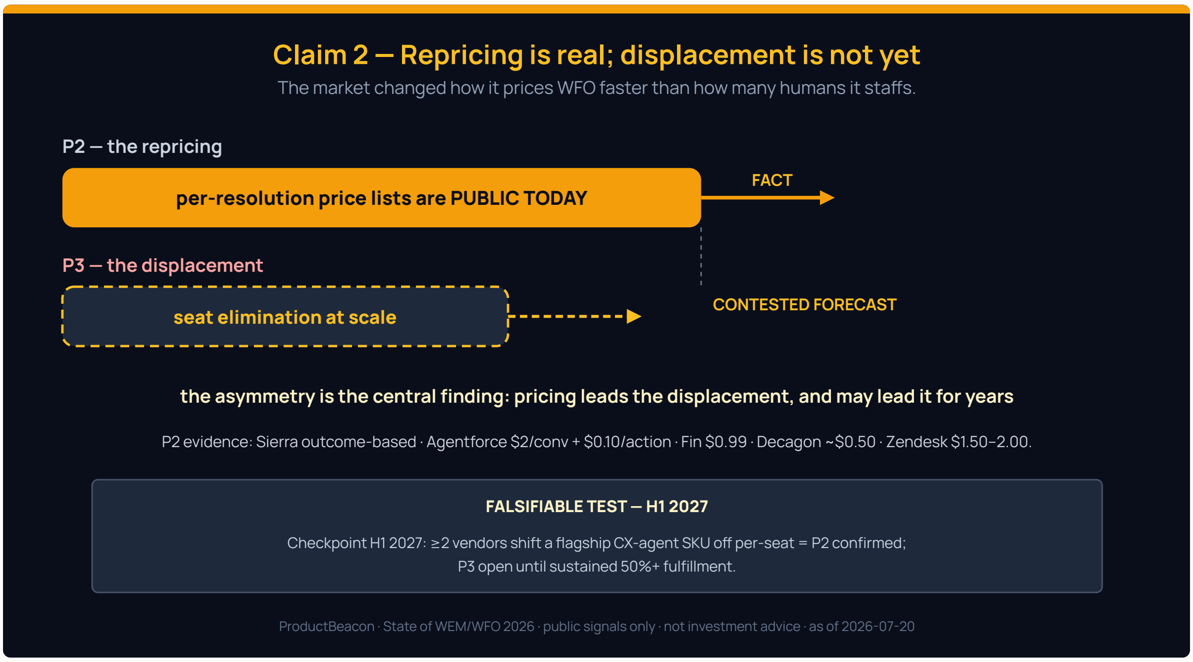

Figure C.6b — Pattern Claim 2 (P2 confirmed / P3 contested): the repricing is an observable fact today; the displacement is a contested forecast.

Pattern Claim 3 — The Unified Falsification Condition (the single falsifier)

Observation. The bet can fail three independent ways, which this synthesis consolidates into **one** dated falsification test rather than three adjacent ones. **My read.** The single-engine-collapse + seat-reset thesis is **downgraded from "confirmed" to "directional, not realized"** if three conditions hold together by **H1 2027**. Condition (a): independent / best-of-breed WEM (the Thoma Bravo Verint+Calabrio platform sold standalone, CallMiner, Observe.AI, the autonomous-natives) is still winning the **majority** of net-new competitive displacements against suite-bundled WEM. **AND** condition (b): **no ≥2** Gravity-tier vendors have publicly shifted a flagship WEM/CX-agent SKU off per-seat to per-outcome pricing. **AND** condition ©: no named enterprise reports **sustained** autonomous fulfillment above **50%** on a primary channel without a public reversal. If all three hold, the collapse is a marketing frame; if any two fail, the thesis stands. *Honesty note: clause (a) is the analyst-judgment leg (there is no public dataset of net-new WEM displacements), whereas (b) and © are publicly observable. The single cleanest real-time tell remains any Gravity vendor reverting off per-resolution to per-seat.* **Conditional prediction.** As above: **H1 2027** is the dated checkpoint; the three clauses are the threshold. The single most important real-time signal to watch: **any Gravity vendor publicly reverting off per-resolution back to per-seat** (the cleanest single falsifier of P2, the proxy currently carrying the thesis). **Sources.** The report's original falsification condition from the scope definition; the fulfillment-reversal form in the Autonomous CX Winning & Losing section, reconciled here into one falsifier.

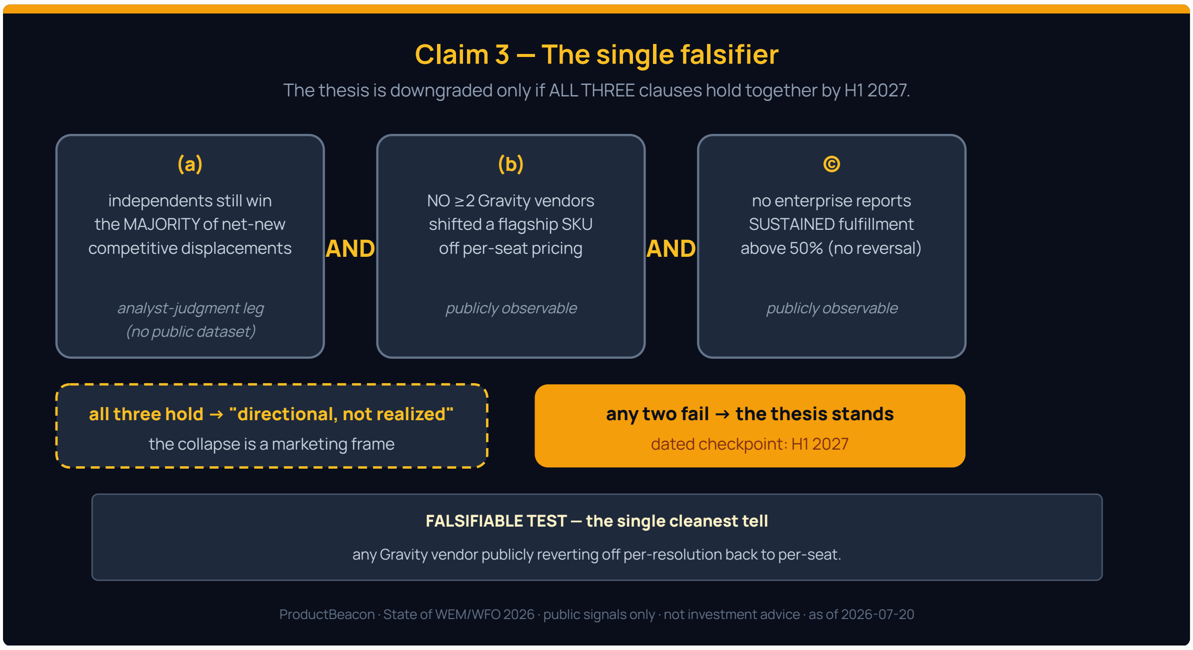

Figure C.6c — Pattern Claim 3: the unified falsification condition — three clauses dated H1 2027, plus the single cleanest real-time tell.

Winners

Reading the per-market verdicts together, this is my opinion across the whole stack, not a scoreboard. **The suites: NICE, Genesys, Salesforce, Five9.** They win the "one platform, already integrated, one engine" pull that recurs in the Quality, Performance, Assist, and Autonomous verdicts; the collapse rewards whoever the buyer already runs. **Thoma Bravo CX (Verint + Calabrio).** It is the private-equity-consolidated platform large enough to be its own suite: the wildcard that is neither pure independent nor legacy incumbent.

Watch

The autonomous-native independents: Sierra, Decagon, Parloa. This is the one market where independent capital genuinely contests the category rather than a feature. Whether frontier funding out-runs the installed base is the open question. **The best-of-breed holders.** CallMiner and Observe.AI in Quality, Assembled and Peopleware in WFM, AmplifAI in Performance, and Balto in Assist hold on depth, un-owned data, and outcome measurement. **The unified falsification condition.** The three clauses are dated H1 2027 in Pattern Claim 3. The single cleanest tell is any Gravity vendor reverting off per-resolution back to per-seat.

Losers

No graded vendor carries a public distress signal in any market's Funding & Fault-Lines section. What is losing, across markets, is a *structural position*, and the evidence is restructuring rather than failure: the standalone best-of-breed WEM business is folding into the consolidation engines. Verint was taken private ($2.0B, closed 26 Nov 2025) and merged with Calabrio. Cogito was absorbed into that platform. NICE bought Cognigy (~$955M) and LiveVox (~$350M) into its owned stack. Read specific to each vendor's WEM/WFO business across markets, these are consolidation events, not deaths. The one position the whole report marks as genuinely exposed is per-seat licensing itself. The per-resolution price lists of Proxy P2 are already displacing it, in public.

The Naming Migration Tracks the Collapse

WFO → WEM → "CX-AI" is not just a label drift; it tracks the collapse. **"Workforce Optimization"** named an era when the product was *modules a human-staffed contact center bought* (record, score, schedule, coach). **"Workforce Engagement Management"** renamed it when the pitch moved to *the agent's experience and retention*. The 2026 vendor language is the third rename in progress: NICE's "the era of AI **and** human agents," the "Agentic Experience Engine," "manage performance across human **and** AI agents on the same queues." The workforce being optimized, engaged, and managed is **no longer assumed to be human.** The category renamed itself each time the unit of work changed, and it is renaming itself again now. The freshness window (to 20 Jul 2026) put the third rename on new surfaces: Microsoft shipped "workforce engagement management" natively in Dynamics 365, to "plan, operate, and improve **human and AI workforce** performance" (microsoft.com, 22 Jun 2026). Salesforce launched Agentforce Contact Center WEM (salesforce.com, 22 Jun 2026). Zendesk's pricing page now carries a "Workforce

Engagement Bundle" SKU (zendesk.com/pricing, page updated 16 Jul 2026). The unit-of-work rename is spreading beyond the incumbent set. That is why a report that leads with "WEM" and carries "WFO" is documenting a live migration, not choosing a synonym.

The verdict

- **P1 (single-engine): Confirmed at the product/SKU surface.** Strongest at NICE/Five9/Salesforce; visible as the control-surface gate. It is not a claim about how the engine is built under the hood.
- **P2 (seat → outcome): Confirmed as observable pricing fact.** Per-resolution price lists are public today. This is the cleanest evidence in the report.
- **P3 (displacement): Open / contested.** Fulfillment is rising, but seat elimination at scale is a forecast with a live counter-signal (Klarna reversal). This is where the bet can still be lost.

Net: the State of WEM/WFO 2026 finds a market that has **already collapsed its products onto one engine and re-priced itself per outcome.** It has done so **ahead of — and not yet proven by — the human-seat displacement that the whole story is ultimately about.** The repricing is the present tense; the displacement is the bet. The falsification condition (in the Winning & Losing section) is dated H1 2027.

The per-market Winning/Losing verdicts consolidate into a best-of-breed-vs-suite read, market by market. Suites lead on the "one platform, already integrated, one engine" pull (Quality, Performance, Assist, Autonomous). Independents hold on depth, the un-owned data, and outcome-measurement: CallMiner/Observe.AI in Quality, Assembled/Peopleware in WFM, AmplifAI in Performance, Balto in Assist, and Sierra/Decagon/Parloa in Autonomous. The Thoma Bravo Verint+Calabrio platform is the wildcard: a PE-consolidated "independent" large enough to be its own suite.

---

1. Five9, "Five9 Launches Agentic CX with AI Agents that can Reason, Decide, and Take Action" (BusinessWire, 10 Jun 2025) + five9.com/products/capabilities/ai-agents — the "Agentic Experience Engine" named as powering both Agent Assist and the autonomous AI Agents. Vendor pillar, quoted as said. Accessed 2026-06-16. https://www.businesswire.com/news/home/20250610324767/en/Five9-Launches-Agentic-CX-with-AI-Agents-that-can-Reason-Decide-and-Take-Action and https://www.five9.com/products/capabilities/ai-agents ↩
2. NICE, CXone Mpower Autopilot / Mpower materials (nice.com) — single-engine-gated-SKU positioning ("support employees via Mpower Copilot or serve customers directly via Mpower Autopilot, tapping into the full CXone Mpower ecosystem"). Vendor positioning, quoted as said. Accessed 2026-06-16. https://www.nice.com/announcements/2025-customer-contact-center-award-nice-cxone-mpower-autopilot-powers-smart-self-service-cx ↩
3. Sierra, "Outcome-based pricing for AI Agents" (billed when the agent achieves a successful resolution; no charge on escalation). Vendor pricing disclosure. Accessed 2026-06-16. https://sierra.ai/blog/outcome-based-pricing-for-ai-agents ↩
4. Saksoft, "The Evolution of Salesforce Agentforce Pricing" ($2/conversation; Flex Credits $0.10/action) + Jitendra Zaa, "Salesforce Agentforce Credits & Cost Model: Complete Guide 2026." Vendor pricing disclosure (named-outlet). Accessed 2026-06-16. https://www.saksoft.com/post/the-evolution-of-salesforce-agentforce-pricing-

what-changed-and-what-it-means-for-you/ and https://www.jitendrazaa.com/blog/salesforce/salesforce-agentforce-credits-cost-model-complete-guide-2026/ ↩

5. Fin by Intercom, "Fin AI Agent Resolutions" ($0.99 per resolution/outcome; $49 base; 50 included monthly resolutions) + Fin AI, "AI Agent Pricing Comparison 2026." Vendor pricing disclosure. Accessed 2026-06-16. https://fin.ai/help/en/articles/10772642-fin-ai-agent-resolutions and https://fin.ai/learn/ai-customer-service-agent-pricing-comparison ↩

6. Decagon, "What is resolution-based pricing?" (per-resolution model; reported ~$0.50/resolution on a $50K platform fee; no public pricing page) + eesel AI, "How much does Decagon AI actually cost? A 2026 Decagon pricing breakdown." Reported pricing. Accessed 2026-06-16. https://decagon.ai/glossary/what-is-resolution-based-pricing and https://www.eesel.ai/blog/decagon-pricing ↩

7. Fin AI, "AI Agent Pricing Comparison 2026: Cost Guide" (Zendesk $1.50 committed / $2.00 pay-as-you-go per automated resolution). Named-outlet multi-vendor pricing comparison. Accessed 2026-06-16. https://fin.ai/learn/ai-customer-service-agent-pricing-comparison ↩

8. NICE, "NiCE Launches CXone Mpower Agents: Enterprise-Grade Agentic AI Agents Built for CX to Deliver Automated Fulfillment" (BusinessWire, 17 Jun 2025) + NICE CXone Mpower Autopilot (nice.com) — "AI that gets things done." Vendor pillar, quoted as said. Accessed 2026-06-16. https://www.businesswire.com/news/home/20250617282648/en/NiCE-Launches-CXone-Mpower-Agents-Enterprise-Grade-Agentic-AI-Agents-Built-for-CX-to-Deliver-Automated-Fulfillment and https://www.nice.com/announcements/2025-customer-contact-center-award-nice-cxone-mpower-autopilot-powers-smart-self-service-cx ↩

9. Genesys, "Agentic Virtual Agent" (genesys.com/capabilities/virtual-agents) + "Genesys Launches AI Agents with Greater Autonomy" (BusinessWire, 9 Sep 2025) — "Conversations don't resolve customer issues: execution." Vendor pillar, quoted as said. Accessed 2026-06-16. https://www.genesys.com/capabilities/virtual-agents and https://www.businesswire.com/news/home/20250909495245/en/Genesys-Launches-AI-Agents-with-Greater-Autonomy-to-Drive-Enterprise-Wide-Customer-Experience-Orchestration ↩

10. Lorikeet, homepage (lorikeetcx.ai) + "Lorikeet Raises $35M Series A" (PRNewswire, Aug 2025) — "whether it's a three-step form or a 40-step refund flow, Lorikeet executes with precision." Vendor pillar, quoted as said. Accessed 2026-06-16. https://www.lorikeetcx.ai/ and https://www.prnewswire.com/news-releases/lorikeet-raises-35-million-usd-series-a-to-deliver-ai-powered-universal-customer-concierges-able-to-solve-real-challenges-not-just-chatbots-302522821.html ↩

11. Salesforce, "Lessons from 500,000 Agentforce Customer Conversations" (Agentforce resolving 84%+ of questions across ~380,000 conversations on help.salesforce.com), corroborated by CX Today, "Salesforce Resolves 85% of Its Customer Queries with Agentforce." Vendor-reported, labeled. Accessed 2026-06-16. https://www.salesforce.com/news/stories/agentforce-customer-support-lessons-learned/ and https://www.cxtoday.com/crm/salesforce-resolves-85-of-its-customer-queries-with-agentforce/ ↩

12. TechCrunch, "Sierra raises $950M as the race to own enterprise AI gets serious" (Series E $950M at ~$15.8B, 4 May 2026, GV + Tiger Global; prior Series C $350M at ~$10B, 4 Sep 2025). Private company; named-outlet is the primary-tier source. Accessed 2026-06-16. https://techcrunch.com/2026/05/04/sierra-raises-950m-as-the-race-to-own-enterprise-ai-gets-serious/ ↩

13. CMSWire, "Decagon triples valuation to $4.5B with $250M Series D" (closed 28 Jan 2026, Coatue + Index) + Decagon Series C announcement ($131M at ~$1.5B, 23 Jun 2025). Prior/current marks stated distinctly. Accessed 2026-06-16. https://www.cmswire.com/contact-center/decagon-triples-valuation-to-45b-with-250m-series-d/ and https://decagon.ai/resources/series-c-announcement ↩

14. Parloa newsroom, "Parloa valued at $3 billion with $350M Series D" (15 Jan 2026, General Catalyst) + Parloa Series C ($120M at ~$1B, 6 May 2025). Accessed 2026-06-16. https://www.parloa.com/parloa-in-the-press/parloa-valued-at-3-billion-with-350m-series-d/ and https://www.parloa.com/parloa-in-the-press/parloa-raises-120m-series-c-to-reinvent-customer-service-with-agentic-ai/ ↩

15. Customer Experience Dive, “Klarna reinvests human talent in customer service AI chatbot” (CEO Siemiatkowski: cut “too far, too fast”; humans reintroduced for complex cases). Named-outlet, reported as an attributed claim. Accessed 2026-06-16. https://www.customerexperiencedive.com/news/klarna-reinvests-human-talent-customer-service-AI-chatbot/747586/ ↩

16. Gartner, “Gartner Predicts Agentic AI Will Autonomously Resolve 80% of Common Customer Service Issues Without Human Intervention by 2029” (2025-03-05). Directional prediction. Accessed 2026-06-16. https://www.gartner.com/en/newsroom/press-releases/2025-03-05-gartner-predicts-agentic-ai-will-autonomously-resolve-80-percent-of-common-customer-service-issues-without-human-intervention-by-20290 ↩

17. Gartner, “Hype Cycle for Agentic AI” — prediction that over 40% of agentic-AI projects will be canceled by end of 2027 (cost, unclear value, inadequate risk controls). Counter-signal. Accessed 2026-06-16. https://www.gartner.com/en/articles/hype-cycle-for-agentic-ai ↩

18. NICE, “NiCE to Acquire Cognigy” (announced 28 Jul 2025, ~$955M) + SEC Form 6-K (closed 8 Sep 2025). Accessed 2026-06-16. https://www.nice.com/press-releases/nice-to-acquire-cognigy-advancing-the-leading-cx-ai-platform-to-accelerate-ai-first-customer-experience and https://www.sec.gov/Archives/edgar/data/0001003935/000117891325002467/exhibit_99-1.htm ↩

19. CX Today, “Playvox Confirms NICE Takeover with a Name Change” (Playvox close 1 Oct 2024, terms undisclosed). Accessed 2026-06-16. https://www.cxtoday.com/workforce-engagement-management/playvox-confirms-nice-takeover-with-a-name-change/ ↩

20. Thoma Bravo, “Thoma Bravo Completes Acquisition of Verint, a Leader in AI-Driven Customer Experience Automation” (take-private closed 26 Nov 2025), PRNewswire; and Verint press room, “Verint Agrees to Be Acquired by Thoma Bravo for $2 Billion” (agreed 25 Aug 2025, $20.50/share). Accessed 2026-06-16. https://www.prnewswire.com/news-releases/thoma-bravo-completes-acquisition-of-verint-a-leader-in-ai-driven-customer-experience-automation-302636053.html and https://www.verint.com/press-room/2025/verint-agrees-to-be-acquired-by-thoma-bravo-for-2-billion/ ↩

21. Channel Futures, “In Major CCaaS Play, NICE to Acquire LiveVox for $350 Million” (Oct 2023, closed Dec 2023). Accessed 2026-06-16. https://www.channelfutures.com/mergers-acquisitions/in-major-ccaas-play-nice-to-acquire-livevox-for-350-million ↩

Disclosures

DISCLOSURE.

This is ProductBeacon's independent, impartial market research, based solely on publicly verifiable signals, published on the open web. No vendor sponsors, no paywalled data, no analyst-firm reuse. It reflects ProductBeacon's independent view and is not investment advice. Full independence and sourcing methodology: Methodology →.

NOT ADVICE.

This report does not constitute investment, legal, tax, or accounting advice. No claim should be relied upon as the basis for any investment decision. The author holds no trading position in any named public security and is not compensated by any named vendor. Readers who use this report in investment contexts bear sole responsibility for their decisions.