PRODUCTBEACON RESEARCH · WEM RADAR

Three vendors put a number on the AI line.

Weekly edition, 12 August 2026. New since 5 August. A busy week, and it cut both ways.

By Yohay Etsion · former NICE product leader (Recording/WFO), 2008–2021 · public sources only · full disclosure on the last slide

THE WINDOW

Seven days. One read.

- This edition covers **5 August to 12 August 2026**, a seven-day window, against the **State of Workforce Engagement Management 2026** report published on 29 July.
- The universe is the report's graded companies across its five markets: Quality and Interaction Analytics, Workforce Management, Performance and Engagement, Agent Assist, and Autonomous Customer Experience.
- This is a weekly edition, so there is **no Pattern-Claims scoreboard**. The first full scoreboard is due between 19 and 26 August.

> Public-source search only. LinkedIn was not collected and no securities-filing feeder ran. Every read below is our interpretation of a public, dated signal, never asserted as fact.

THE AI LINE

Three vendors, three denominators.

- **Five9** reported on **6 August** that AI revenue grew **78% year over year** and now accounts for **15% of subscription revenue**, up from 9% a year earlier.
- **NiCE** put AI annual recurring revenue at **15% of cloud revenue**.
- **Genesys**, in its annual sustainability report published **11 August**, disclosed that about **20% of new business annual contract value** came from Genesys Cloud AI.
- Read the denominators carefully, because they differ. Five9's share is of *subscription* revenue, NiCE's is of *cloud* revenue, and the Genesys figure is of *new business contract value*, not of revenue. They are not the same measurement.

> As we read it, what these three share is that each now considers the AI line material enough to disclose separately. That, rather than any one number, is the finding.

THE PRICING STRUCTURE

Money that moves between humans and agents.

- Alongside its print, **Five9** confirmed a shift toward a **revenue-commitment model** rather than pure per-seat licensing. Customers commit to a spend level over three to five years, then reallocate that spend between human agent seats and AI agents as the mix changes.
- Its roughly **$100 million total contract value** win with a Fortune 100 financial services firm, closed through the Google Cloud Marketplace, was structured this way.
- This is **not per-resolution pricing**. It is a capacity commitment, which is the model **Bain & Company**'s June 2026 analysis of about 200 business-to-business software companies found roughly four in five vendors actually choose when they move beyond seats.

> For the report's claim that the workforce planning unit becomes blended human and AI, this is the first time the radar has seen it written into a gravity-tier vendor's commercial structure rather than its marketing.

THE SUITES

The largest deal, and the caveat under it.

- **NiCE** disclosed its **largest-ever CXone and Cognigy deal**: a nine-digit total contract value agreement, described as eight-digit annual contract value, with **HM Revenue & Customs**, the United Kingdom tax authority, delivered with **Capgemini**. A national tax authority is a useful data point against the idea that agentic AI stays confined to low-risk service journeys.
- NiCE reported **AI annual recurring revenue up 52% year over year**, at 15% of cloud revenue. Read it precisely: **52% is growth on the AI annual recurring revenue line, not on total revenue**, and **15% is a share of cloud revenue, not of total revenue**. These are the vendor's own figures.
- The caveat came in the question-and-answer session. An analyst noted that **net-new AI annual recurring revenue looked lower than the year-ago quarter despite record AI bookings**. Chief financial officer Beth Gaspich attributed this to "the conversion of timing".

> As we read it, record bookings with softer net-new annual recurring revenue is the shape of a market where the contract signs well before the deployment lands. Signed commitments are not yet displaced seats.

THE OTHER HALF

The displacement evidence got weaker.

- **TaskUs AI Services growth slowed to 25.8%.** That is its first quarter below 30% in six quarters, with margin narrowing to 18.7%. TaskUs is the outsourcer furthest into the AI transition.
- **Gartner found only 20% of organizations have cut agent headcount because of AI.** That reading is as of late 2025. Roughly 80% plan to move staff into new roles instead.
- **NiCE net-new AI annual recurring revenue came in below the year-ago quarter.** Record bookings, softer net-new. The detail is on the previous slide.

> As we read it, the repricing keeps accruing evidence while the seat displacement stays contested. That asymmetry is the finding.

A CORRECTION WE OWE

We told you half of it last week.

- Last week we said the European Union's AI Omnibus did not move the prohibitions or the Article 50 transparency obligations. That is true. It was also only half the picture.
- **The high-risk obligations moved by 16 months.** The Digital Omnibus on AI, in force 27 July 2026, deferred Annex III stand-alone high-risk systems from 2 August 2026 to 2 December 2027.
- **Annex III covers this market directly.** It lists monitoring or evaluating workers' performance and behaviour, which is the workforce engagement surface almost exactly.
- **The emotion rule did not move.** Article 5(1)(f) prohibits inferring worker emotions from biometric/voice data in the workplace, in force since February 2025, with a medical/safety carve-out as the only exception. It attaches to the deploying employer's use, not the vendor's sale.

> As we read it, the Emotion-AI Wall claim is unaffected and arguably sharpened. The workplace emotion prohibition is live now. The compliance machinery around workforce-monitoring AI is 16 months out.

THE QUIET PART

Three consecutive silent weeks.

WORKFORCE MANAGEMENT

Assembled, **Peopleware**, **Aspect**, **Centrical**, **Cisne**. No tracked pure-play made a public move this window.

QUALITY AND INTERACTION ANALYTICS

Cresta, **Level AI**, **CallMiner**, **OttoQA**.

PERFORMANCE AND ENGAGEMENT

AmplifAI, **Spinify**, **Zenarate**.

AUTONOMOUS CUSTOMER EXPERIENCE

Sierra, **Decagon**, **Parloa**, **Lorikeet**.

ALSO QUIET

Sprinklr, **Uniphore**, **Observe.AI**. **Thoma Bravo CX** (Verint and Calabrio) made no in-window announcement.

> Three consecutive quiet weeks for the independent middle, while the gravity tier prints earnings and quantifies its AI line, is the shape the report predicted for the thinning-middle claim. It is not yet evidence of absorption. It is evidence of a widening disclosure gap between vendors who publish numbers and vendors who do not. Quiet here also means quiet in public-source search: LinkedIn was not collected.

THE RECEIPTS

Every claim here is public and dated.

The sources behind the moves in this carousel:

COMPANY AND PRIMARY ANNOUNCEMENTS

Five9 · Genesys · Puzzel · European Commission

BUSINESS AND TECHNOLOGY PRESS

CX Today · CX Foundation · No Jitter · CMSWire · TechTarget · StockTitan · TradingView · Yahoo Finance

LEGAL, REGULATORY AND RESEARCH ANALYSIS, QUOTED AND ATTRIBUTED

Gibson Dunn · Freshfields · Ogletree Deakins · KPMG Law · Lexology · JD Supra · Bain & Company · Gartner

Full citations, claim by claim and link checked, available on request.

FULL DISCLOSURE

No claim broke this window.

On the evidence this search surfaced, no Pattern Claim was falsified and no counter-tell fired. We saw no gravity vendor revert from per-resolution to per-seat pricing, which the report holds as the single cleanest counter-tell. Evidence accrued on both sides. The first full scoreboard is due between 19 and 26 August.

→ Read the full report

productbeacon.agency/research/state-of-wfo-2026

> The author was employed by NICE from 9 April 2008 to 20 May 2021, including senior product-leadership responsibility for the Recording / Workforce-Optimization line of business, and remains bound by a confidentiality and trade-secret undertaking from that employment. The report relies solely on published material and contains no information derived from that employment. NICE is graded on the identical published-material rubric applied to every peer, no softer, no harsher, and was held under an adversarial cross-chapter read in both directions. The author holds no NICE equity and has no current commercial relationship with NICE.